

Retrieval Engine Competition Project

Notebook Report

Group: SeaFour

Maksym DOLHOV

Mehdi AGHAEI

Nguyen Ho Bao KHANH

Report objective. This version updates the written report to precisely match our experimental progression and the current saved notebook. We detail our journey from TF-IDF and BM25+ baselines to a robust dense embedding retriever. The active configuration evaluates and submits only the embedding retriever, combined with a TF-IDF + LinearSVC category classifier and a dominant-category post-filter.

Primary source files. `kaggle-submission.ipynb`, `solutions_SeaFour.csv`, and this L^AT_EX report.

Notebook file	<code>kaggle-submission.ipynb</code>
Submission file	<code>solutions_SeaFour.csv</code>
Current default model	<code>embedding</code>
Submission size target	Top 7,500 retrieved documents before category filtering
Date	March 19, 2026

1 Notebook-Aligned System Overview

The logic is organized into several stages: strict data loading, shared text normalization, sequential model experimentation (TF-IDF to BM25+ to Embeddings), aggressive caching for performance, offline evaluation, and submission generation.

The current system relies entirely on a highly optimized dense first-stage retriever paired with a dominant-category filter.




Figure 1: Active end-to-end workflow in the saved notebook.

Component	Implemented	Active in saved run	Role in the pipeline
TF-IDF retrieval	Yes	No	Sparse lexical baseline using cosine similarity. Fast, but struggled with semantic mismatch.
BM25+ retrieval	Yes	No	Stronger lexical baseline with token-level scoring and document-length normalization.
Embedding retrieval	Yes	Yes	Dense semantic first-stage retrieval with Sentence-Transformers. Yielded the best metrics.
Category classifier	Yes	Yes	Predicts one category per query for evaluation and final CSV output.
Dominant-category filter	Yes	Yes	Keeps only documents belonging to the dominant category among the top 20 retrieved items.

Table 1: What the notebook supports versus what the saved execution actually uses.

2 Dataset and Preprocessing

The notebook loads 216,041 documents, 327 training queries, 141 test queries, and one training qrels file. The document collection is distributed across five categories, with significant class imbalance.




Figure 2: Document distribution by category.

Text normalization is shared across all components: separators `-`, `_`, and `/` are replaced by spaces, whitespace is collapsed, and text is lowercased. Retrieval documents are built from `title + text + tags`, while retrieval queries use only `title + text`. Average normalized lengths are 900.9 characters for documents and roughly 48.5 for queries. This length disparity informed our ultimate shift away from purely lexical methods.

3 Experimental Progression and Pipeline Optimization

Our methodology evolved through a structured series of experiments, beginning with simple baselines and moving toward a complex, optimized semantic architecture.

3.1 Phase 1 & 2: Lexical Baselines (TF-IDF and BM25+)

We initially implemented a **TF-IDF** retriever (using unigrams/bigrams with `min_df=2`). This served as a fast, lightweight baseline to validate our data loading and evaluation pipeline. While transparent, TF-IDF struggled significantly with semantic mismatch when queries paraphrased document concepts without sharing exact vocabulary, recall dropped.

To improve this, we upgraded to **BM25+**. By tokenizing the text and applying `k1=1.5` and `b=0.75`, BM25+ provided much stronger lexical ranking, correcting for the high variance in document lengths. However, BM25+ still fundamentally relied on exact token overlap, meaning synonymous phrasing remained a weakness.

3.2 Phase 3: The Semantic Shift (Dense Embeddings)

To overcome the vocabulary mismatch problem, we transitioned to dense semantic retrieval using the Sentence-Transformer model `all-MiniLM-L6-v2`. Documents and queries were encoded into 384-dimensional L2-normalized vectors, allowing us to retrieve documents based on contextual meaning rather than exact keywords. This dramatically improved recall.

3.3 Phase 4: Overcoming the 1.5-Hour Bottleneck

The shift to embeddings introduced a severe computational roadblock: initially, encoding the 216,041 documents and calculating scores took almost **1.5 hours to run**. This latency crippled our ability to run rapid iterations.

To solve this, we implemented aggressive caching and chunking mechanisms:

- **Disk Caching:** We engineered a caching system that fingerprints the dataframes and saves the expensive document embeddings as `.numpy` arrays. Repeated notebook runs now load the matrix in seconds instead of re-encoding.
- **Chunked Query Scoring:** Instead of creating massive, memory-crashing score blocks, we scored queries in chunks of 32 (`EMBEDDING_QUERY_CHUNK_SIZE = 32`).
- **Partial Sorting:** We replaced full array sorts with `np.argpartition` to extract only the top- K candidates efficiently.

3.4 Phase 5: Hyperparameter Tuning and Category Filtering

With the pipeline running smoothly, we began tuning our parameters. We experimented heavily with `top_k` retrieval depths. Because our dataset features a high semantic mismatch, we found that selecting a very large initial candidate pool (`top_k = 7,500`) safely maximized recall.

To refine this massive list, we trained a TF-IDF + LinearSVC Category Classifier. By looking at the top 20 retrieved documents, we identify a "dominant category" and filter out any subsequent documents that do not belong to that category. This successfully cleans the noise out of the long tail of the 7,500 candidates.

4 Evaluation Protocol and Results

The notebook evaluates the optimized embedding run against the 327 training queries. Because the dominant-category filter may return fewer than K documents, precision is calculated based on the number of predicted documents actually present after filtering. The combined score is calculated as:

$$\text{Score} = \frac{1}{4} (\text{Recall} + \text{Precision} + \text{MRR} + \text{Accuracy})$$

Model	TopK	Recall	Precision	MRR	Accuracy	Score
Embedding	7,500	0.91572	0.00178	0.46989	0.92661	0.59957

Table 2: Offline metrics reported by the saved notebook execution.

The recall (0.9157) remains highly robust even after the category filter is applied. The filter effectively trims the fatretained documents per query drop to a mean of roughly 4,841, without leaving any query empty. The MRR (0.46989) confirms that despite the long list, the first relevant document surfaces quite early in the ranking.

5 Conclusion

Our system evolved from basic lexical matching to a highly optimized semantic retrieval pipeline. By implementing dense embeddings, we solved the semantic mismatch problem; by engineering a robust caching and chunking system, we solved the 1.5-hour computational bottleneck; and by applying a dominant-category post-filter, we successfully cleaned the noise from deep `top_k` retrievals. This architecture represents the final, submitted SeaFour pipeline.